

Advanced methods in Machine Learning

Assignment 3

Isak Falk (16095283)

Github: isakfalk

Due on 11th of April, 2017

April 11, 2017

Contents

1 Problem A: Cart-pole	3
1.1 Generate three trajectories under a uniform random policy. Report the episode lengths and the return from the initial state.	3
1.2 Generate 100 episodes under the above random policy, and report the mean and standard deviation of the episode lengths and the return from the starting state.	4
1.3 Collect 2000 episodes under a uniformly random policy and implement batch Q-learning to learn to control the cart-pole.	4
1.3.1 Linear	4
1.3.2 100-hidden units MLP	6
1.4 Function approximation using Q-learning.	8
1.5 Network with ReLUs and different hidden units.	9
1.6 Experience replay.	11
1.7 Compare using a target network	12
1.8 Compare Sarsa to previous performance	13
2 Problem B: Atari games	14
2.1 Report the score and frame counts from the three games under a random policy, evaluated on 100 episodes. Report both the average and the standard deviation.	14
2.2 Report performance on the three games from an initialized but untrained Q-network, evaluated on 100 episodes. Explain why the performance can be different from part one.	14
2.3 Plot the losses of your agent during the course of training. Discuss the shapes of the curves and why they might differ from supervised learning.	14
2.3.1 Pong	15
2.3.2 Ms Pacman	16
2.3.3 Boxing	17
2.4 Report the final performance in terms of cumulative undiscounted rewards per episode of your trained agent, averaged over 100 episodes. Describe any modifications to the above by which you were able to improve performance in this limited data regime.	18

1 Problem A: Cart-pole

Note on value function representation For all of the value functions we have used the form where we take as input the current state s_t which then maps through Q to the action space, giving a score to each possible action. For the cart pole we have that $a_t \in \{0, 1\} = A$ and $s_t \in \{([-2.4, 2.4], \mathbb{R}, [-41.8, 41.8], \mathbb{R})\} = S$, which represents cart position, cart velocity, pole angle and pole velocity at tip. This means that we have a value function $Q : S \rightarrow A$. In order to find the action-value for the pair (s_t, a_t) we just look at index a_t of $Q(s_t)$ such that $Q(s_t, a_t) := [Q(s_t)]_{a_t}$.

How statistics was collected All of the statistics was collected by bringing the agent out of the training regime and evaluated on the environment using the learned greedy policy on 20 episodes normally, although this differs slightly depending on how computationally demanding the task was and what was asked for in the question.

I saved the models every time I saw an improvement in test performance while running training, which is why most of them get close to 300 episode length when evaluated.

Hyperparameters For each of the questions we used similar hyperparameters. For optimization we used ADAM since it worked better than SGD and similar or better than RMSprop even though RMSprop is said to be suitable for RL problems. The networks did not use extra implementations such as regularisation and/or batch normalisation since this didn't seem necessary.

I used the following learning rate, batch size, buffer size and epochs for each subquestion.

Question	optimizer	decay schedule	learning rate	batch size	buffer size	epochs/episodes	Hidden units
A3	ADAM	None	Various	16	-	500	-
A4	ADAM	None	0.0001	-	-	2000	100
A5	ADAM	None	0.0001	-	-	2000	30/1000
A6	ADAM	Polynomial	0.0001	32	10000	2000	100
A7	ADAM	None	0.0001	32	10000	2000	100

For the polynomial decay I used the builtin Tensorflow `tf.train.polynomial_decay` such that the learning rate that you choose is the starting learning rate, the end learning rate was set to be 0.000001, decay steps 100000 and power to 0.5.

1.1 Generate three trajectories under a uniform random policy. Report the episode lengths and the return from the initial state.

We have the following episode lengths and total discounted returns from the initial state under a uniformly random policy:

Run	1	2	3
Episode length	20	19	16
Total discounted reward	-0.8179	-0.8262	-0.8515

1.2 Generate 100 episodes under the above random policy, and report the mean and standard deviation of the episode lengths and the return from the starting state.

We get the following figures:

	mean	std.
Episode length	22.53	11.717
Total discounted reward	-0.8027	0.0898

1.3 Collect 2000 episodes under a uniformly random policy and implement batch Q-learning to learn to control the cart-pole.

We first sample 2000 episodes worth of tuples of $(S_t, A_t, R_{t+1}, S_{t+1})$, through a uniformly random policy. After gathering about 44000 such tuples, as the mean length of an episode is about 22, we train agents running offline batch Q-learning to learn the optimal policy of controlling the cart-pole. For this part we consider 5 different learning rates, $[10^{-5}, 10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 0.5]$ and two different value function approximators:

Linear

$$Q(S_t, A_t) = [S_t^T W + b]_{A_t}$$

Since $|A| = 2$ and $|S| = 4$ we have that if we represent vectors in column form that $W \in \mathbb{R}^{4 \times 2}, b \in \mathbb{R}^{1 \times 2}$.

MLP For the Multi-Layer Perceptron we have that

$$Q(S_t, A_t) = [\sigma(S_t^T W_1 + b_1)W_2 + b_2]_{A_t}$$

Were we have the freedom to choose the hidden layers H such that $W_1 \in \mathbb{R}^{4 \times H}, b \in \mathbb{R}^H, W_2 \in \mathbb{R}^{H \times 2}, b_2 \in \mathbb{R}^2$, and σ represents an elementwise sigmoid function, in our case the ReLU. This is true for all the different networks below as well, with different H depending on the subquestion. In this case $H = 100$.

1.3.1 Linear

We have the following plots

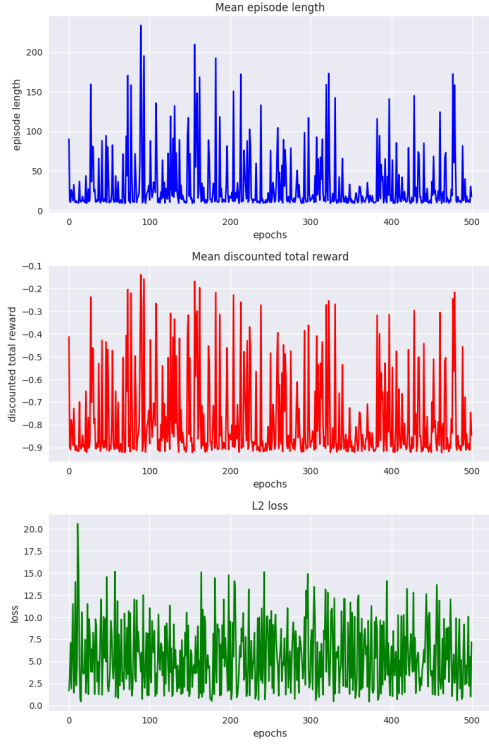


Figure 1: Learning rate: 0.5

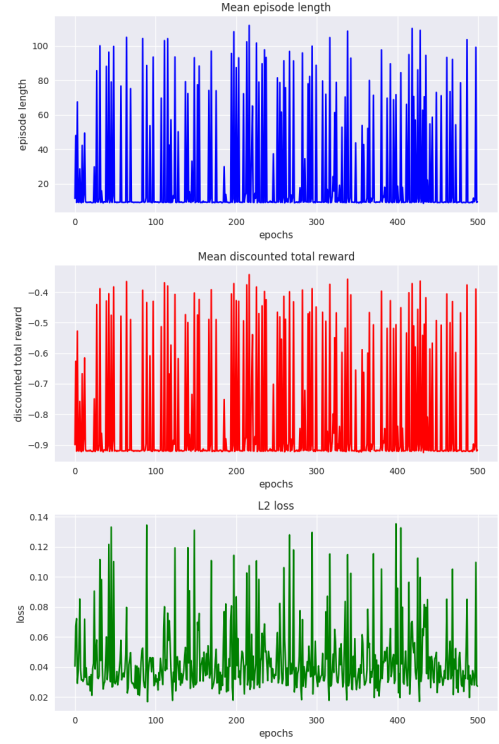


Figure 2: Learning rate: 0.1

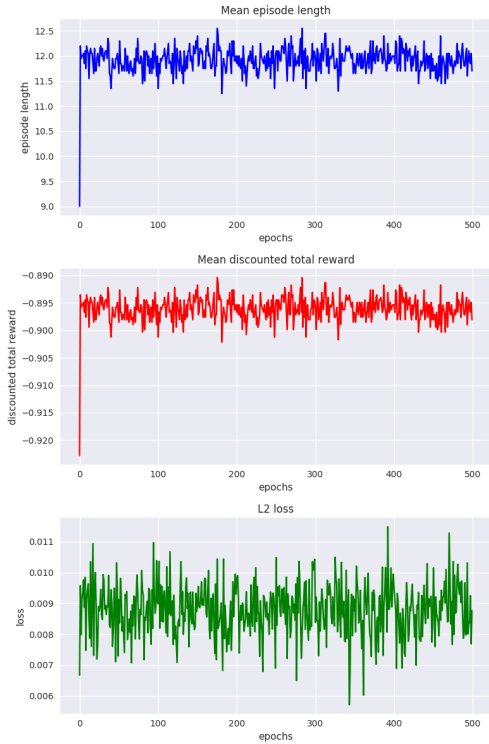


Figure 3: Learning rate: 0.01

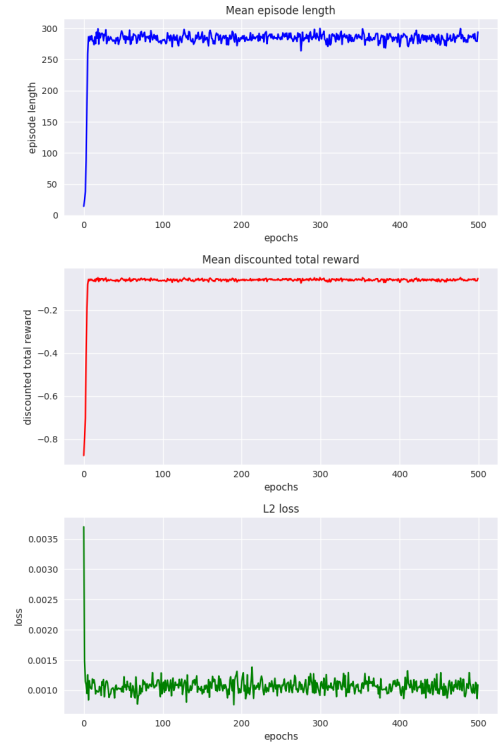


Figure 4: Learning rate: 0.001

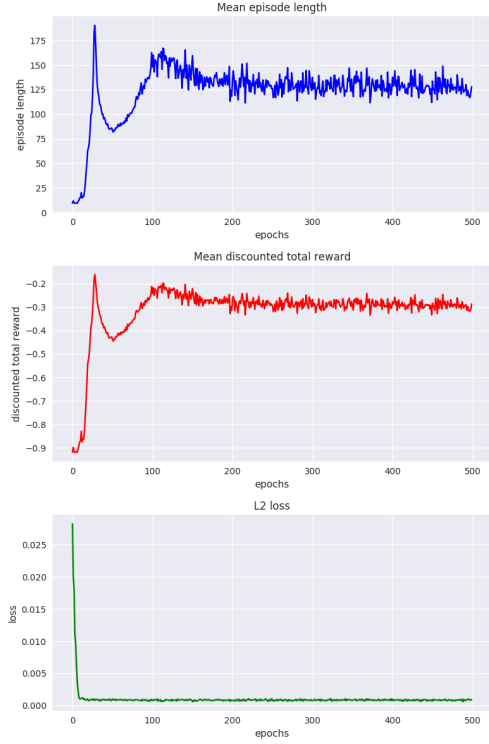


Figure 5: Learning rate: 0.0001

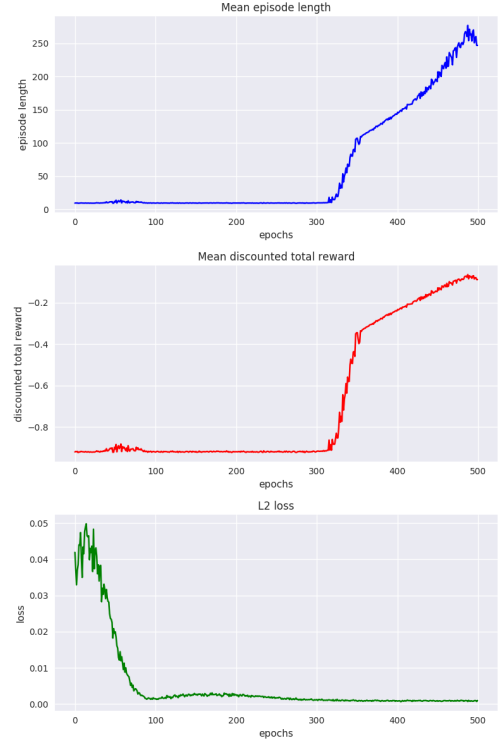


Figure 6: Learning rate: 0.00001

1.3.2 100-hidden units MLP

We have the following plots

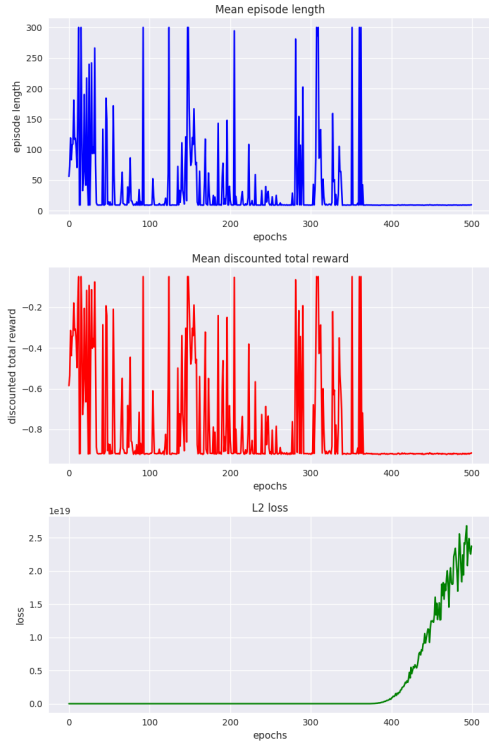


Figure 7: Learning rate: 0.5

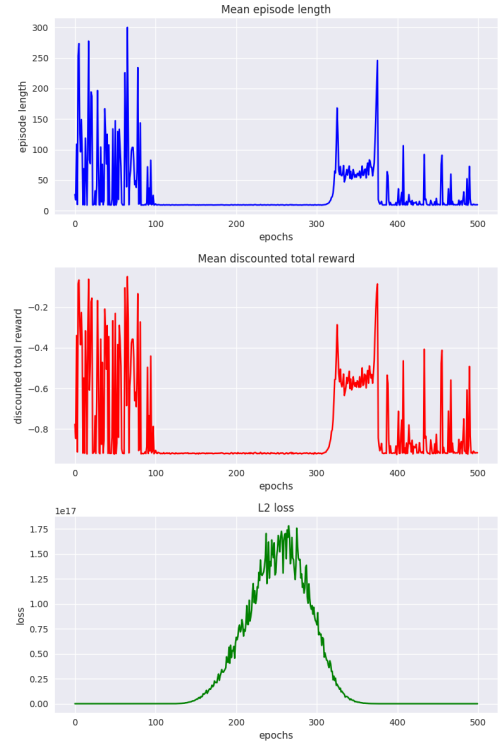


Figure 8: Learning rate: 0.1

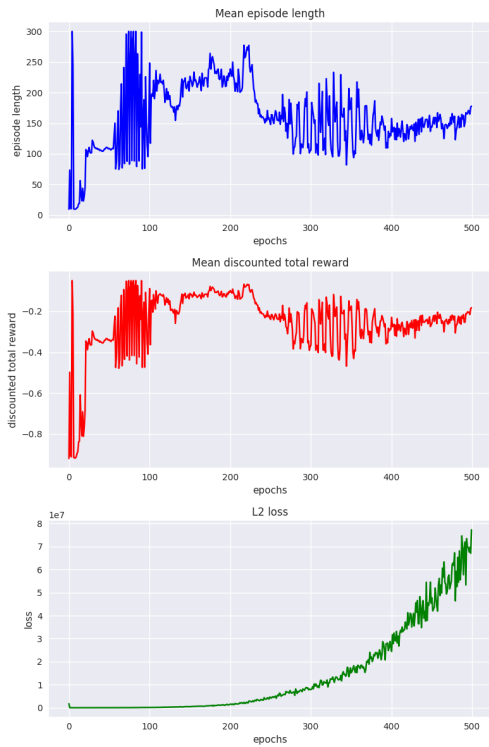


Figure 9: Learning rate: 0.01

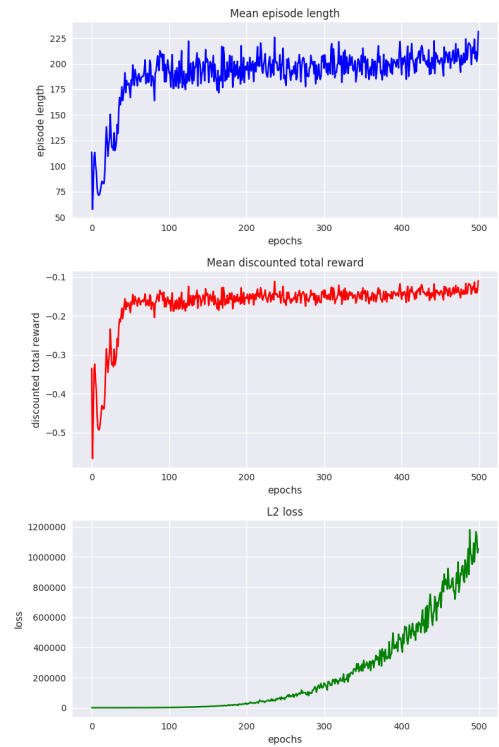


Figure 10: Learning rate: 0.001




Figure 11: Learning rate: 0.0001




Figure 12: Learning rate: 0.00001

From all this we see that the linear learner does a pretty good job at solving the Cartpole problem, getting to 300 and hovering there for a learning rate of 0.001. However, for other values we have more erratic behaviour and given the volatile nature of RL, it might have been pure luck since other runs for lower learning rates give less stable runs as can be seen in the figures for learning rate 0.0001 and 0.00001.

While the neural net is more powerful than the linear learner as it can represent a bigger space of functions, it is also more volatile and prone to catastrophic forgetting. An interesting phenomenon is that in many of these runs the neural networks learn a decent policy (such as in the figure of learning rate 0.0001) even though the bellman loss seems to converge. An example of destructive learning can be found in the figure of the learning rate 0.00001.

1.4 Function approximation using Q-learning.

We implement an online Q-learning algorithm using an MLP for function approximation as specified above. We let the hidden layer be of size 100 and non-linearity being the ReLU function outputting a value for each possible action 0 and 1.

Using this we have that the bellman residual is

$$\delta = R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)$$

and the loss

$$Loss = 0.5\delta^2.$$

We train the agent using an epsilon greedy policy with $\epsilon = 0.05$ over 2000 episodes. To evaluate performance we let the agent run 20 episodes using a greedy policy with regards to the current Q_t at the finished training episode. During evaluation we don't train. We then take the mean of the episode length, total discounted reward and Bellman loss over these 20 evaluation episodes and let the means be representative of the agent at that point in time.

We plot all of the learning traces using a small alpha for the opacity and the line specifies the mean of the 100 traces.




Figure 13: Learning traces

From the plot we see the inherent volatility and variance using the same algorithm (Q-learning) on the same environment (CartPole-v0) with identical hyperparameters. From the traces we see that at most episodes during training we have traces which have an average episode length of 9 to 300, echoing the fact that they will be agents that oscillate between optimal (300) and worst case (9) scenario. However, on average we have decent behaviour of around 200. However, since we cap the maximal episode length at 300, we have that the mean will always be below 300 and thus be biased towards a lower value of episode lengths than if we actually evaluated the episode length of an unconstrained environment.

If we look at the Bellman residue we can see that this shows that we almost always overestimate our current Q-action-value function compared to the target, as found in the Double Q-learning paper¹.

1.5 Network with ReLUs and different hidden units.

We reuse the same network as in the previous question but with different number of hidden units, 30 and 1000.

¹<https://arxiv.org/pdf/1509.06461.pdf>




Figure 14: 30 hidden units




Figure 15: 1000 hidden units

From the graphs we see the volatility of online Q-learning. I found while training that the 30-unit networks converged quicker than the 1000 unit networks, which makes sense since the smaller network should be faster to converge. In both cases, by running and testing for optimal hyperparameters, I experienced that the networks were prone to catastrophic forgetting where they would learn a policy which reach 300 average episode length (Since the evaluation was over 20 episodes, this also means that it must have gotten a length of 300 for all episodes since we take the mean.) essentially solving the problem, only to forget everything later and getting a policy that consistently gave an average episode length of about 9.5.

The difference between the 30 and the 1000 unit network was not as big as I would have thought. However, our problem here consists of finding the optimal network over an online period of 2000 episodes. This means that we are not really finding the network that learns the best policy overall, but the best policy over 2000 episodes (which is a pretty big constraint as we normally train over epochs of data sets of size 2000 and more). I imagine that training over a longer period of episodes and a lower learning rate would mean that the more flexible 1000 hidden unit network would learn a better policy and potentially mitigate the catastrophic forgetting.

1.6 Experience replay.

We reuse the same network as in previous questions but add an experience replay buffer where we store the previous tuples up to a certain point, which is specified by the buffer size. This should stabilise the learning since we can both batch and reuse previous experience to learn.

We have the following graph




Figure 16: 100 hidden units

Compared to the previous question we see that the experience replay buffer has a stabilising effect on the

bellman residual, meaning that we converge to a policy that is consistent much quicker. We also note that the overestimation of the Q-value seem to have disappeared since the Bellman residual is now distributed more evenly around zero and not consistently negative as before.

The experience replay buffer makes sure that we can both batch while training, meaning we may learn faster as we take more aspects of the data into consideration as we run backpropagation, stabilise the backpropagation and also decouples the data points making batches for sequential training steps less correlated, while improving data efficiency².

While the bellman residue and thus the loss is improved the most, it also mitigates the volatility of the performance of the network. While it still have catastrophic forgetting, it doesn't seem to be as volatile as in previous questions. However, as we are only doing one run, it's hard to tell properly.

1.7 Compare using a target network

Similarly to above we have a experience replay buffer but also add a target network. Using a target network means that we should stabilise learning as we decouple the update with the current network, removing possible feedback loops due to the target being dependent on the network at the current timestep. By only copying the network at certain intervals we break this dependency since the correlation dies out quickly as we progress in time.

We have the following graph

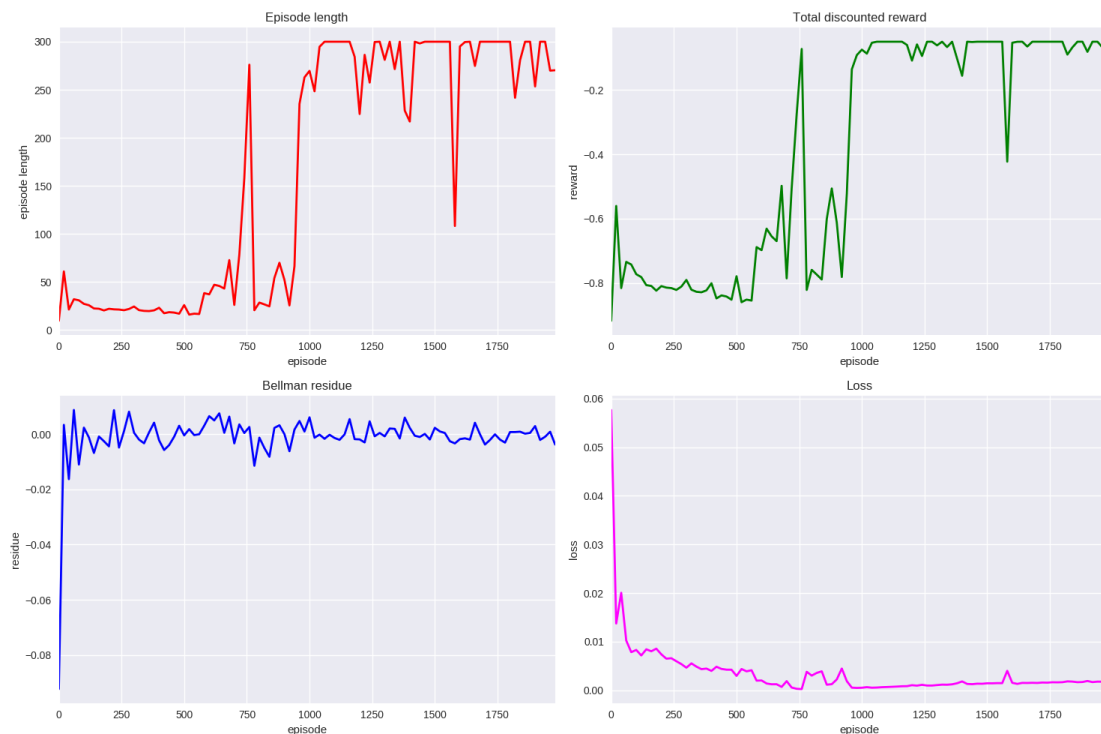


Figure 17: 100 hidden units

It's hard to tell if a target network actually improve things. I think that over time the target network should be able to use the data more efficiently rather than over the period of just 2000 episodes.

²p.7, <https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>

What is clear though is that the target network reduces the bias of the Q-value function. Looking at the residual graph we see that it is much more evenly distributed around 0 and not predominantly on one side of the origin (i.e negative as before, due to Q-learning being biased.).

From this we can clearly see the theoretical advantage of a target network. It decouples the target Q-value from the current policy and thus makes the data less biased towards our current policy, leading to data that is less correlated and should be easier to train on.

1.8 Compare Sarsa to previous performance

Not implemented.

2 Problem B: Atari games

2.1 Report the score and frame counts from the three games under a random policy, evaluated on 100 episodes. Report both the average and the standard deviation.

We have the following table

Game	score (mean)	score (std)	frame count (mean)	frame count (std)
Pong	-0.9059	0.2707	1225.86	137.19
Ms. Pacman	2.4412	0.5758	636.03	100.74
Boxing	-0.4100	1.1064	2378.26	13.87

2.2 Report performance on the three games from an initialized but untrained Q-network, evaluated on 100 episodes. Explain why the performance can be different from part one.

We have the following table

Game	score (mean)	score (std)	frame count (mean)	frame count (std)
Pong	-1.1302	0.0317	1018.65	8.96
Ms. Pacman	2.8629	0.1980	578.83	134.55
Boxing	-3.3297	1.4200	1429.05	263.3219

The statistics differ from that of the uniform policy simply since a CNN which is randomly initialized will probably be biased towards different actions even though it is not trained. For example filters which are randomly initialized will often still be able to find some natural shapes (Similar to Gabor filters) and thus act in some manner which is non-random, although suboptimal.

2.3 Plot the losses of your agent during the course of training. Discuss the shapes of the curves and why they might differ from supervised learning.

We have the following plots

2.3.1 Pong




Figure 18: Loss




Figure 19: Residual

2.3.2 Ms Pacman




Figure 20: Loss




Figure 21: Residual

2.3.3 Boxing




Figure 22: Loss




Figure 23: Residual

We see that the plots of the losses have a slow decrease, although these plots are a running average of the loss and thus smoother than they actually are raw. From A we know that the loss is not completely representative of the actual performance of the agent, as we might find a policy which has low loss, but actually is stuck in a bad state, not getting a very good performance (which might be different from the actual discounted reward which relates directly to our loss).

If we were to compare these loss curves to those of Supervised Learning, the difference would have to do with the fact that in RL we are able to interact with the environment, affecting the future states through our actions. In SL we don't do this, instead we simply try to predict the outcome given the input, but we can't change the behaviour of the environment from this. Specifically, in SL we assume that the data is i.i.d while in RL we don't have any such guarantees since the distribution of the states depends on the policy which itself changes with time.

2.4 Report the final performance in terms of cumulative undiscounted rewards per episode of your trained agent, averaged over 100 episodes. Describe any modifications to the above by which you were able to improve performance in this limited data regime.

We have the following table for the averaged cumulative undiscounted rewards per episode of the trained agents.

	Pong	MsPacman	Boxing
Cumulative undiscounted reward	-21.0	21.0	-2.27

As can be seen this is not very good and for the games Pong and MsPacman I consistently get these scores, 21 and -21. Boxing differs from episode and have a higher variability. The reason why my agent fails to learn is twofold.

Since we downsize the frames to 28x28, we almost lose all of the visual cues needed to play the game. There is still enough information to actually learn a policy, but learning it is non-trivial and not easy. I think what happens is that the agent start off gaining some score, as in pacman, and continue this policy indefinitely. Still, the loss function changes so some kind of learning or at least adjustment of the policy through time occurs.

As we only train for a million steps, this might actually be too little to be able to learn a non-trivial policy. We use both experience replay and target network so this in some way decouples the data and should enable the agent to learn in a more stable manner than without it, but the number of iterations might not be enough to actually establish a policy different from the starting one.

I didn't have time to experiment with hyperparameters but given the time I would try different ways of downsampling that would leave more of the image intact, such as clipping it in an intelligent way such that the 28x28 is a patch rather than a downsized version of the original image. I would also look into using more complex learning rate schemes and see if a higher epsilon for exploring would improve the policy as more of the states would be investigated than is now.